

Class 03: models

蕭佳明
遠東科技大學 助理教授

Machine Learning

- **Everything Old Is New Again**
 - Computer capacity
 - Amount of data available
- **You Are Already Reaping the Benefits of Machine Learning**
 - Netflix account
 - Amazon
- **It's All About the Data**



<https://dzone.com/articles/how-will-deep-learning-change-the-way-we-design-ou>

It's All About the Data

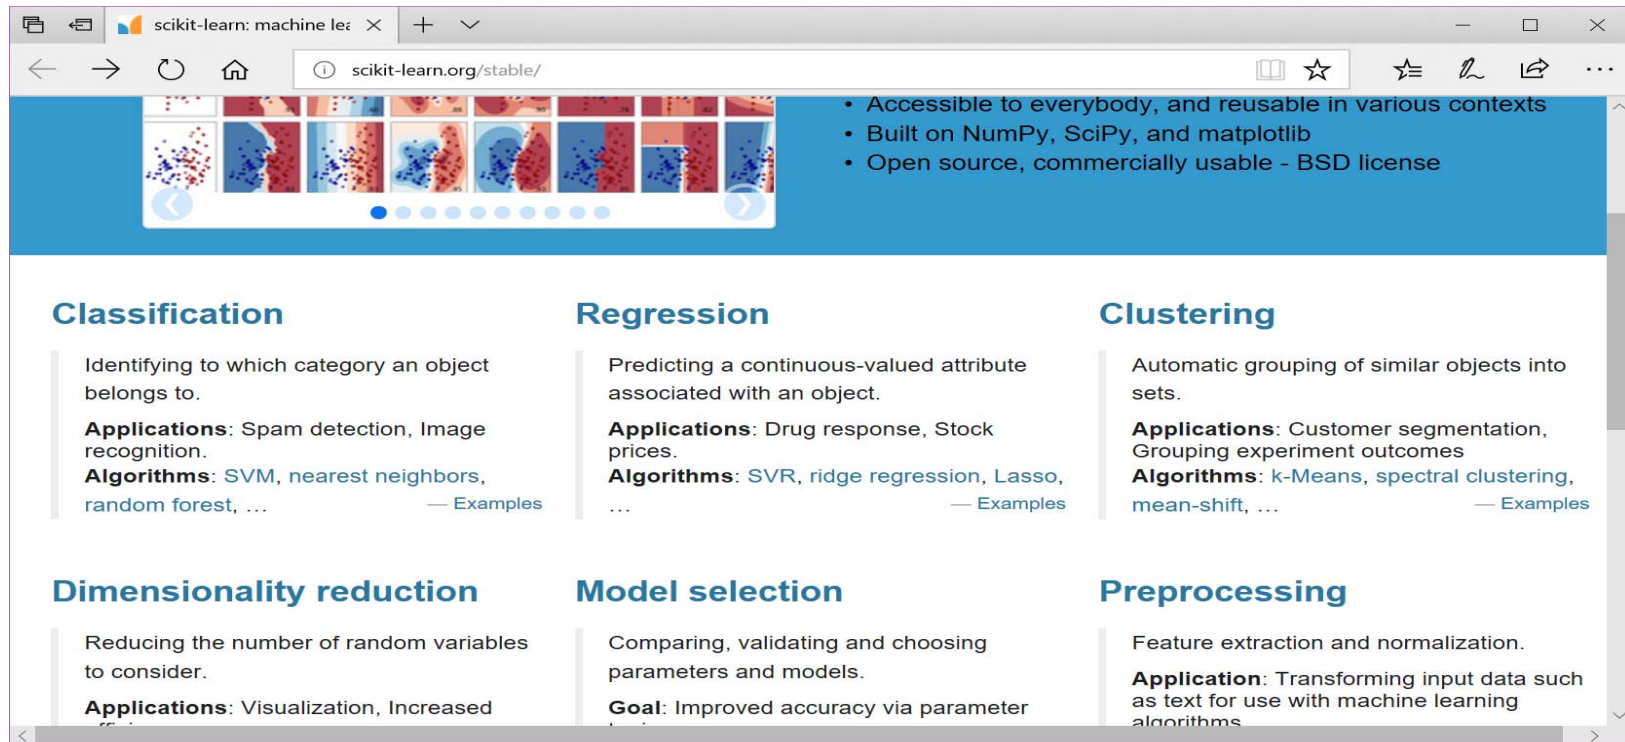
- Given the importance of the data to the success of any machine learning implementation.
 - Data Quality
 - Data Volume
 - Data Timeliness
 - Data Pedigree



<https://dzone.com/articles/how-will-deep-learning-change-the-way-we-design-ou>

Machine Learning

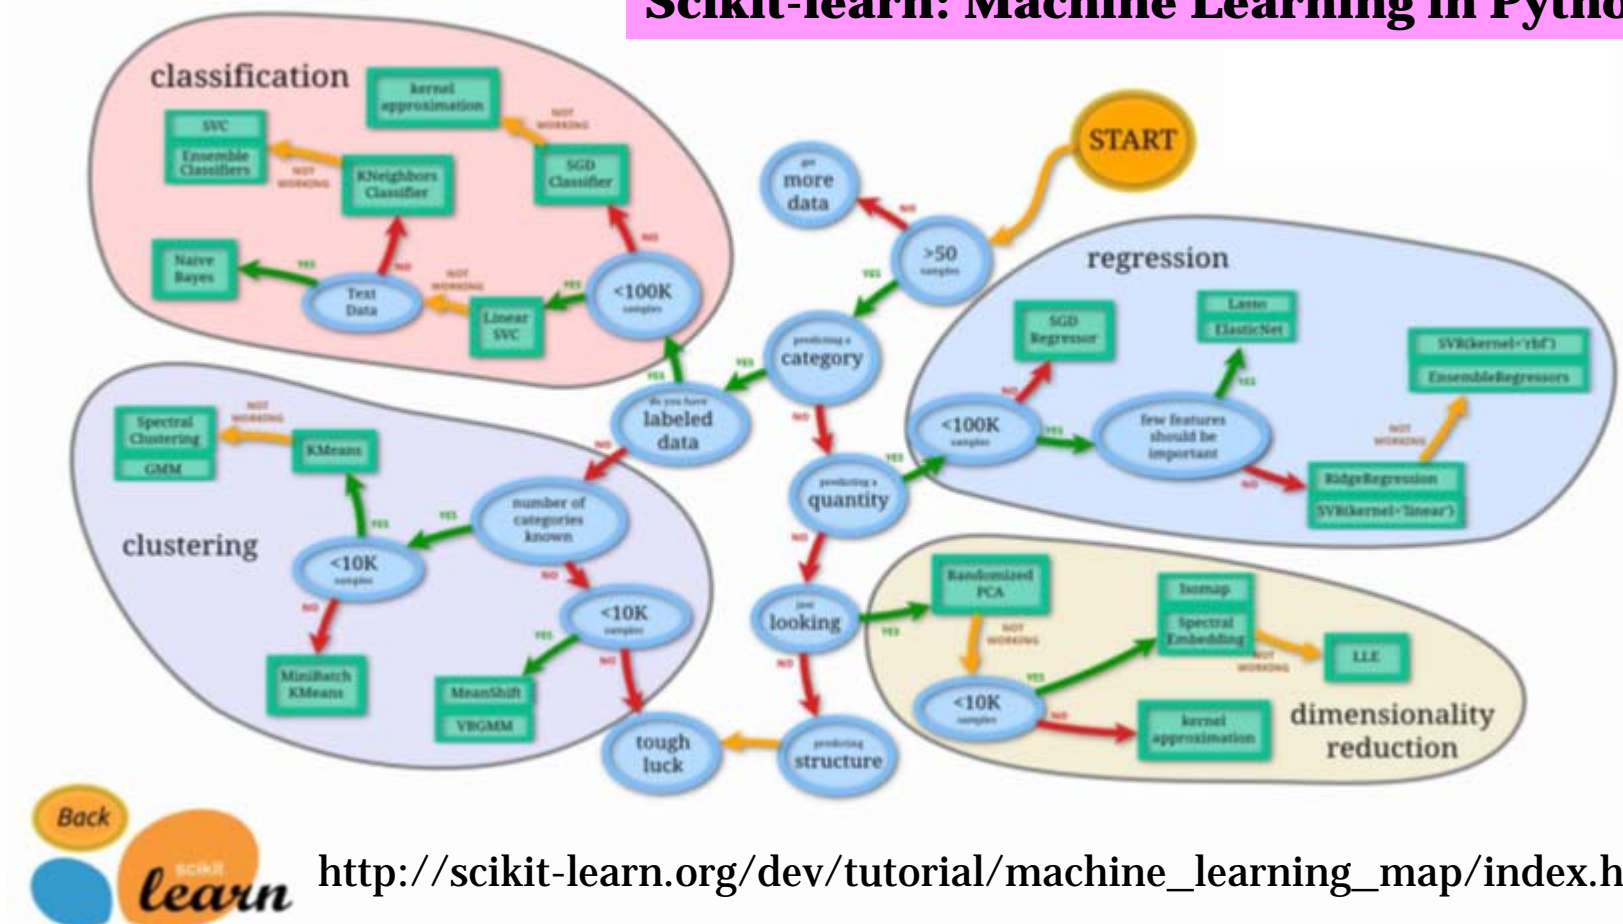
- Scikit-Learn



<http://scikit-learn.org/stable/>

Scikit-Learn Packages

Scikit-learn: Machine Learning in Python



http://scikit-learn.org/dev/tutorial/machine_learning_map/index.html





Supervised Learning

- Model-based learning
 - Linear regression
 - Regression with regularization
 - Logistic regression
 - Support vector machine
 - Decision Tree
 - Random Forests
- Instance-based learning
 - Naive Bayesian model
 - K-nearest neighbor(KNN)

Unsupervised Learning

- Principal Component Analysis
- K-mean

Model Diagnostic




Mapping

$$Y = AX$$

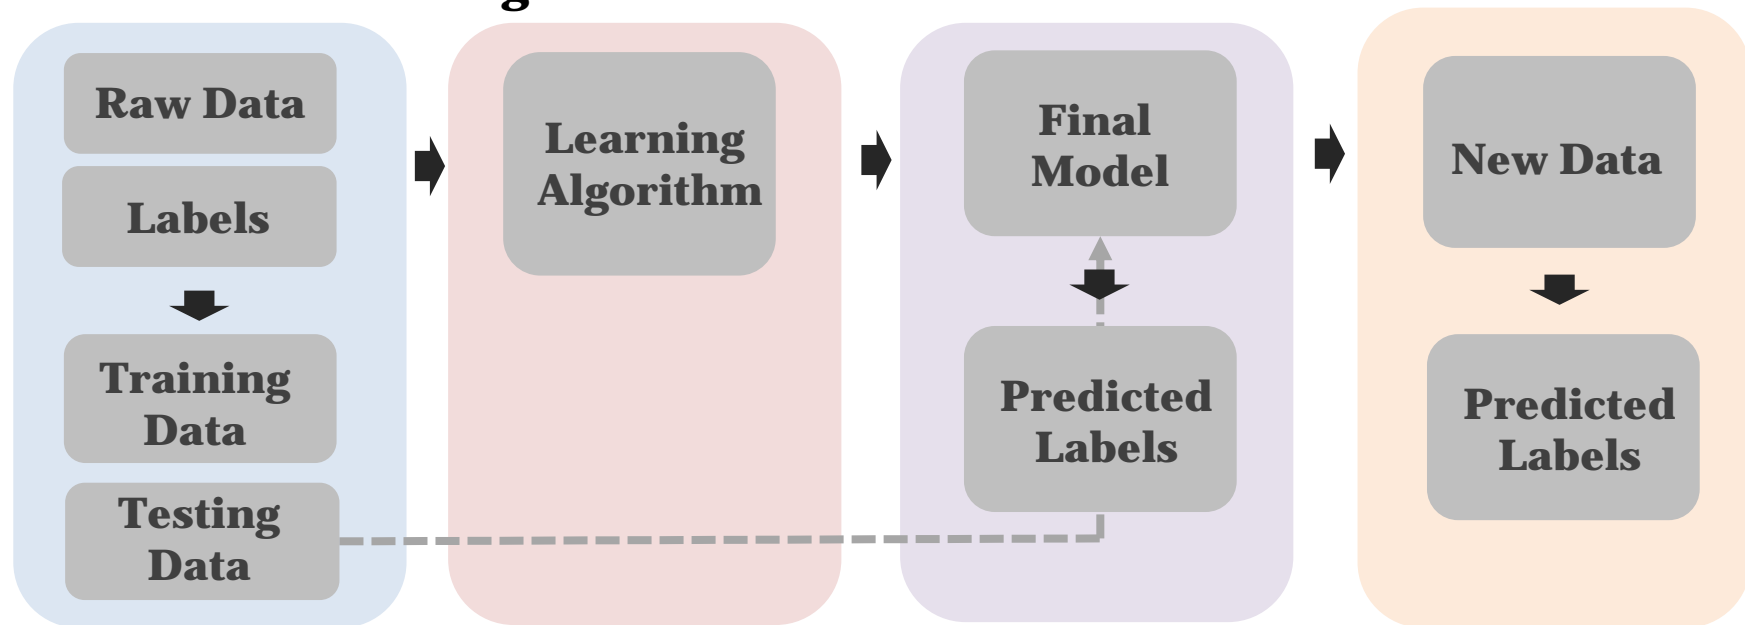
$X: \mathbb{R}^4$

$Y: \mathbb{R}^2$



Supervised Learning

Predictive Modeling



Preprocessing

- Feature extraction -normalization
- Feature selection
- Dimension reduction
- Sampling

Learning


Evaluation

- Model selection
- Cross-validation
- Metric
- Hyperparameter optimization

Prediction

Supervised Learning

- Supervised Learning
 - Model Selection
 - Parameters of Model
 - Parameters of Learning
 - Learning Algorithm
 - Object Function




Toy dataset

- Clustering (data reduction/compression)
- Density estimation

input (x1)	input (x2)
-1	-1
-2	-1
-1	-2
-2	-2
1	1
2	1
1	2
2	2

input (x1)	input (x2)
-1	-1
-2	-1
-1	-2
-2	-2
1	1
2	1
1	2
2	2

input (x1)	input (x2)
-1.5	-1.5
1.5	1.5




Toy dataset

- Data mining

input (x1)	input (x2)
3	1
3	-1
5	-1
5	1
1	1
1	3
-1	1
-1	3


input (x1)	input (x2)
3	1
3	-1
5	-1
5	1
1	1
1	3
-1	1
-1	3

input (x1)	input (x2)
4	0
0	4



Unsupervised Learning

- K-means clustering for data compression



`digits.images.shape = (1797, 8, 8)`

<http://scikit-learn.org/stable/modules/generated/sklearn.cluster.KMeans.html>

06_kmeans_clustering.ipynb



PCA


- Dimensionality reduction
- Data mining

x1	x2
1	3
2	5
3	7
4	9
5	11
6	13
7	15
8	17

PC ₁	PC ₂
-2.02	0.00
-1.44	0.00
-0.86	0.00
-0.28	0.00
0.28	0.00
0.86	0.00
1.44	0.00
2.02	0.00

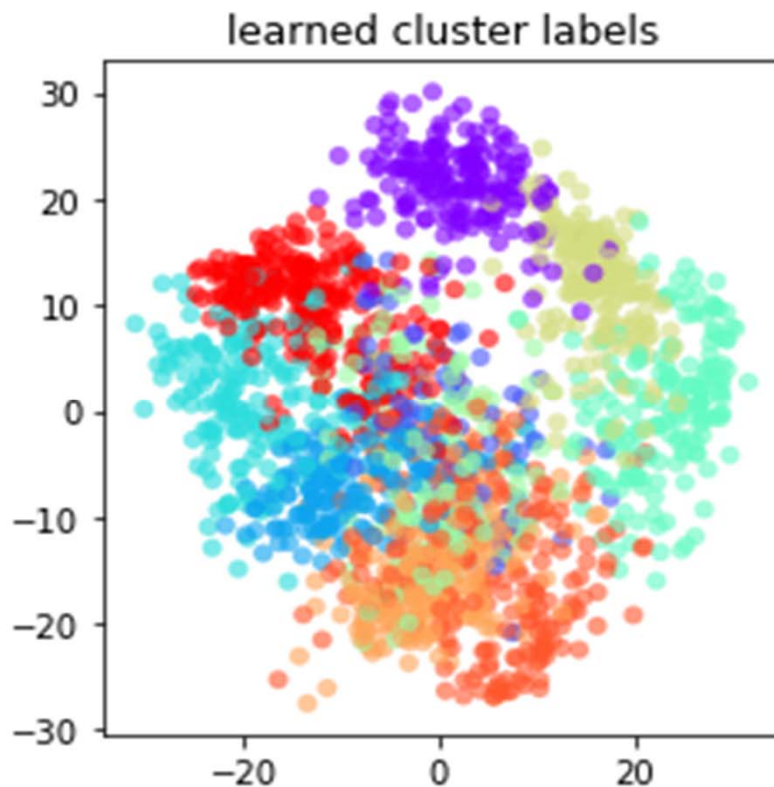
	PC1	PC2
Standard deviation	1.414	1.115e-16

Rotation:		
	PC1	PC2
x	0.7071068	0.7071068
y	0.7071068	-0.7071068




Unsupervised Learning

- PCA Visualization



`digits.images.shape = (1797, 2)`





`digits.images.shape = (1797, 2)`

Reinforcement Learning




- Training data: (S, A, R). (State-Action-Reward)
- Goal: Develop an optimal policy (sequence of decision rules) for the learner so as to **maximize its long-term reward**.




Semi-supervised Learning

- Semi-supervised learning is about using these unlabeled examples to improve supervised learning methods, which generally require labeled examples for training.


Label Spreading 30% data




Label Spreading 50% data



Label Spreading 100% data



SVC with rbf kernel



Active learning with Label Propagation.
Rows show 5 most uncertain labels to learn with the next model.

model 1					
fit with 10 labels	predict: 1 true: 4	predict: 1 true: 8	predict: 1 true: 8	predict: 9 true: 0	predict: 8 true: 3
model 2	4	8	8	0	3
fit with 15 labels	predict: 8 true: 3	predict: 1 true: 4	predict: 8 true: 3	predict: 8 true: 3	predict: 3 true: 3
model 3	3	4	3	3	3
fit with 20 labels	predict: 9 true: 9	predict: 7 true: 1	predict: 8 true: 5	predict: 2 true: 1	predict: 2 true: 7
model 4	9	1	5	2	7
fit with 25 labels	predict: 1 true: 1	predict: 6 true: 6	predict: 1 true: 9	predict: 9 true: 5	predict: 9 true: 5
model 5	1	6	9	5	5
fit with 30 labels	predict: 4 true: 4	predict: 8 true: 8	predict: 3 true: 3	predict: 8 true: 8	predict: 8 true: 8
	4	8	3	8	8

Machine Learning

- **Supervised learning**
- **Unsupervised learning**
- Reinforcement learning
- Semi-supervised learning

Scikit-learn's Estimator

- Available in **all estimators**
 - **model.fit()**: fit training data, X: data, Y: label
 - Supervised: **model.fit(X,Y)** \Leftrightarrow Unsupervised: **model.fit(X)**
- Available in **supervised estimators**
 - **model.predict()** : predict the label of a new set of data by model
 - **model.predict_proba()** : for classification problems, some estimators provides this method to return probability of each class
- Available in **unsupervised estimators**
 - **model.transform()** : transform new data into new basis by model
 - **model.fit_transform()** : some estimators implement this method to efficiently perform a fit and transform on the same input data

Supervised Learning

- In supervised learning, we have a dataset consisting of both **features** (input variables) and **labels** (output variables)
- The task is to construct an **estimator**(model) which enables to predict the labels of an instance given the set of features
- Two categories: **Classification** and **Regression**
 - Classification: the label is **discrete**
 - Regression: the label is **continuous**
- Split into **training** and **testing** datasets

Supervised Learning


- Model-based learning
 - Linear regression
 - Regression with regularization
 - Logistic regression
 - Support vector machine
 - Decision Tree
 - Random Forests
 - Neural Networks
- Instance-based learning
 - Naive Bayesian model
 - K-nearest neighbor(KNN)

Training and Testing Dataset

- To validate the generalization of a trained model, splitting a whole dataset into **training** and **testing** datasets
- We fit and optimize predictive models using the training dataset
- Use testing datasets to evaluate the performance of trained models to ensure the model **generalization**

Generalization

- We hope that a model trained on the basis of a training dataset can seamlessly apply to **unseen testing** dataset
- If the model over fits the training dataset, its performance on testing dataset will be worse
- Higher model complexity, easier to overfitting



Hyperparameters

- Learning the parameters of a prediction function and testing it on **the same data** is a methodological mistake
- There are **Hyperparameters**: choices about the algorithm that we **set** rather than learn

Your Dataset: Images / Corpus

Hyperparameters

- **idea 1:** Choose **hyperparameters** that work best on the data Your Dataset

Your Dataset

- **idea 2:** Split data into train and test, choose **hyperparameters** that work best on test data

Train

Testing

- **idea 3:** Split data into train, val., and test; choose **hyperparameters** on val. and evaluate on test

Train

Validation

Testing

Hyperparameters

- **idea 4: Cross-Validation:** Split data into folds, try each fold as validation and **average** the results

fold1	fold2	fold3	Testing
fold1	fold2	fold3	Testing
fold1	fold2	fold3	Testing

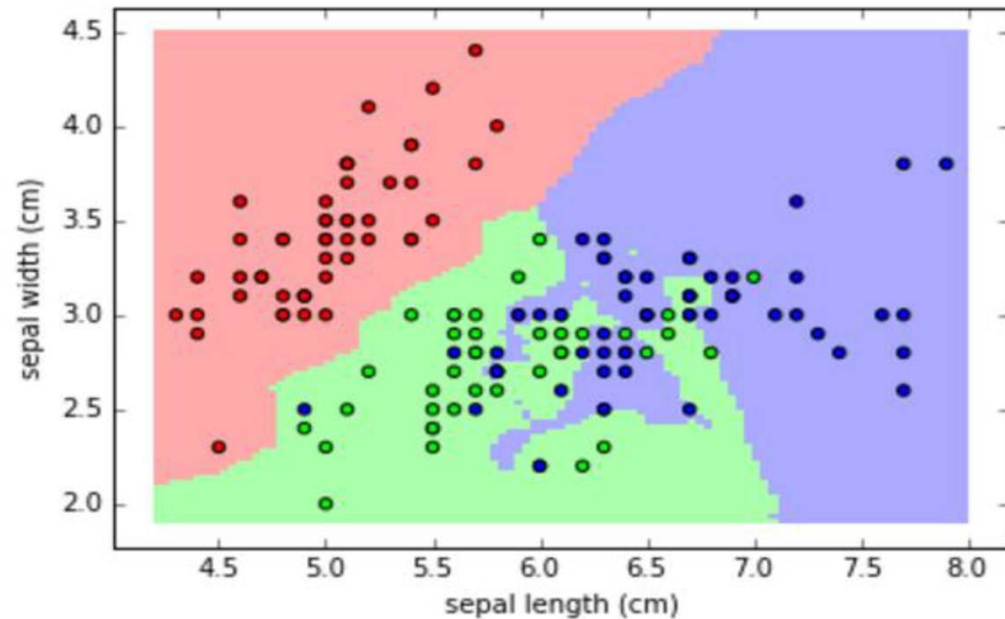
Useful for small datasets, but not used too frequently in deep learning

Classification

- **k-Nearest Neighbor** for iris

Fisher's Iris Data

Sepal length ⇅	Sepal width ⇅	Petal length ⇅	Petal width ⇅	Species ⇅
5.1	3.8	1.9	0.4	<i>I. setosa</i>
4.8	3.0	1.4	0.3	<i>I. setosa</i>
5.1	3.8	1.6	0.2	<i>I. setosa</i>
4.6	3.2	1.4	0.2	<i>I. setosa</i>
5.3	3.7	1.5	0.2	<i>I. setosa</i>
5.0	3.3	1.4	0.2	<i>I. setosa</i>
7.0	3.2	4.7	1.4	<i>I. versicolor</i>
6.4	3.2	4.5	1.5	<i>I. versicolor</i>
6.9	3.1	4.9	1.5	<i>I. versicolor</i>
5.5	2.3	4.0	1.3	<i>I. versicolor</i>
6.5	2.8	4.6	1.5	<i>I. versicolor</i>
5.1	2.5	3.0	1.1	<i>I. versicolor</i>
5.7	2.8	4.1	1.3	<i>I. versicolor</i>
6.3	3.3	6.0	2.5	<i>I. virginica</i>
5.8	2.7	5.1	1.9	<i>I. virginica</i>
7.1	3.0	5.9	2.1	<i>I. virginica</i>
6.3	2.9	5.6	1.8	<i>I. virginica</i>
6.5	3.0	5.8	2.2	<i>I. virginica</i>
7.6	3.0	6.6	2.1	<i>I. virginica</i>
4.9	2.5	4.5	1.7	<i>I. virginica</i>



<http://scikit-learn.org/stable/modules/generated/sklearn.neighbors.NearestNeighbors.html#sklearn.neighbors.NearestNeighbors>

Confusion matrix

	Predicted 0	Predicted 1
Actual 0	True Negative (TN)	False Positive (FP)
Actual 1	False Negative (FN)	True Positive (TP)

$\text{Recall} = \text{TP} / (\text{FN} + \text{TP})$

$\text{Precision} = \text{TP} / (\text{FP} + \text{TP})$

$\text{F1 score} = 2 * (\text{Recall} * \text{Precision}) / (\text{Recall} + \text{Precision})$

Confusion matrix

	Predicted 0	Predicted 1
Actual 0	True Negative (TN)	False Positive (FP)
Actual 1	False Negative (FN)	True Positive (TP)

$$\text{Accuracy} = (\text{TN} + \text{TP}) / (\text{TN} + \text{FP} + \text{FN} + \text{TP})$$

Metrics in Classification

```
from sklearn import metrics
```

```
metrics.confusion_matrix(y_test,y_pred)
```

```
array([[ 7,  0,  0],  
       [ 0,  8,  4],  
       [ 0,  1, 10]])
```

```
metrics.accuracy_score(y_test,y_pred)
```

```
0.8333333333333333
```

```
metrics.precision_score(y_test,y_pred,average="weighted")
```

```
0.8507936507936507
```

```
metrics.recall_score(y_test,y_pred,average="weighted")
```

```
0.8333333333333333
```

```
metrics.f1_score(y_pred=y_pred,y_true=y_test,average="weighted")
```


```
0.8314285714285715
```

```
print(metrics.classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	7
1	0.89	0.67	0.76	12
2	0.71	0.91	0.80	11
avg / total	0.85	0.83	0.83	30

Regression

- Fit a line to the data



Metrics in Regression

- Mean Absolute Error (MAE)
- Mean Square Error (MSE)
- R Squared
- **Peak Signal to Noise Ratio (PNSR)**
- **Structural Similarity Index Measure (SSIM)**

```
metrics.r2_score(y_pred=y_pred,y_true=y_test)
```

```
0.82031173595813334
```

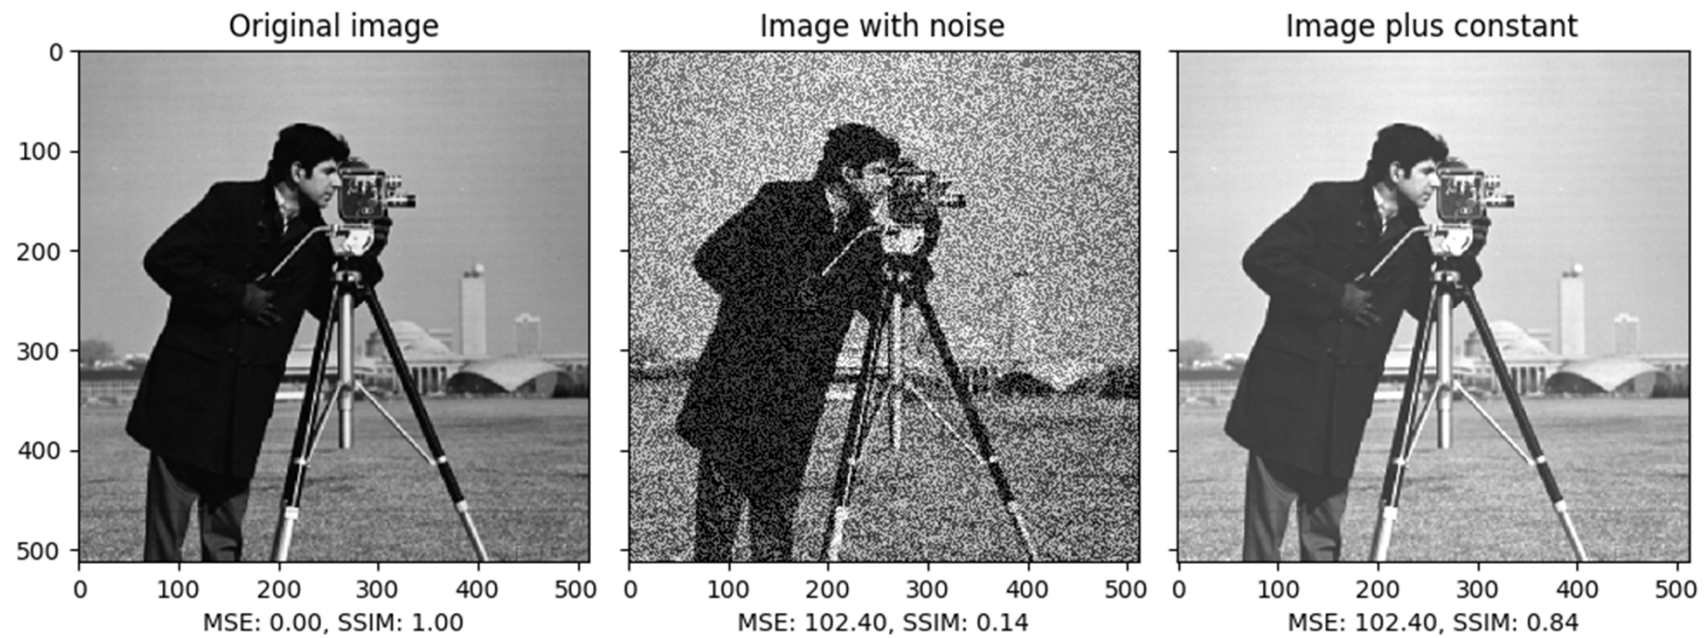
```
metrics.mean_absolute_error(y_pred=y_pred,y_true=y_test)
```

```
0.67860317120632041
```


```
metrics.mean_squared_error(y_pred=y_pred,y_true=y_test)
```

```
0.62402155001086912
```

SSIM




Data Scientist's Toolbox



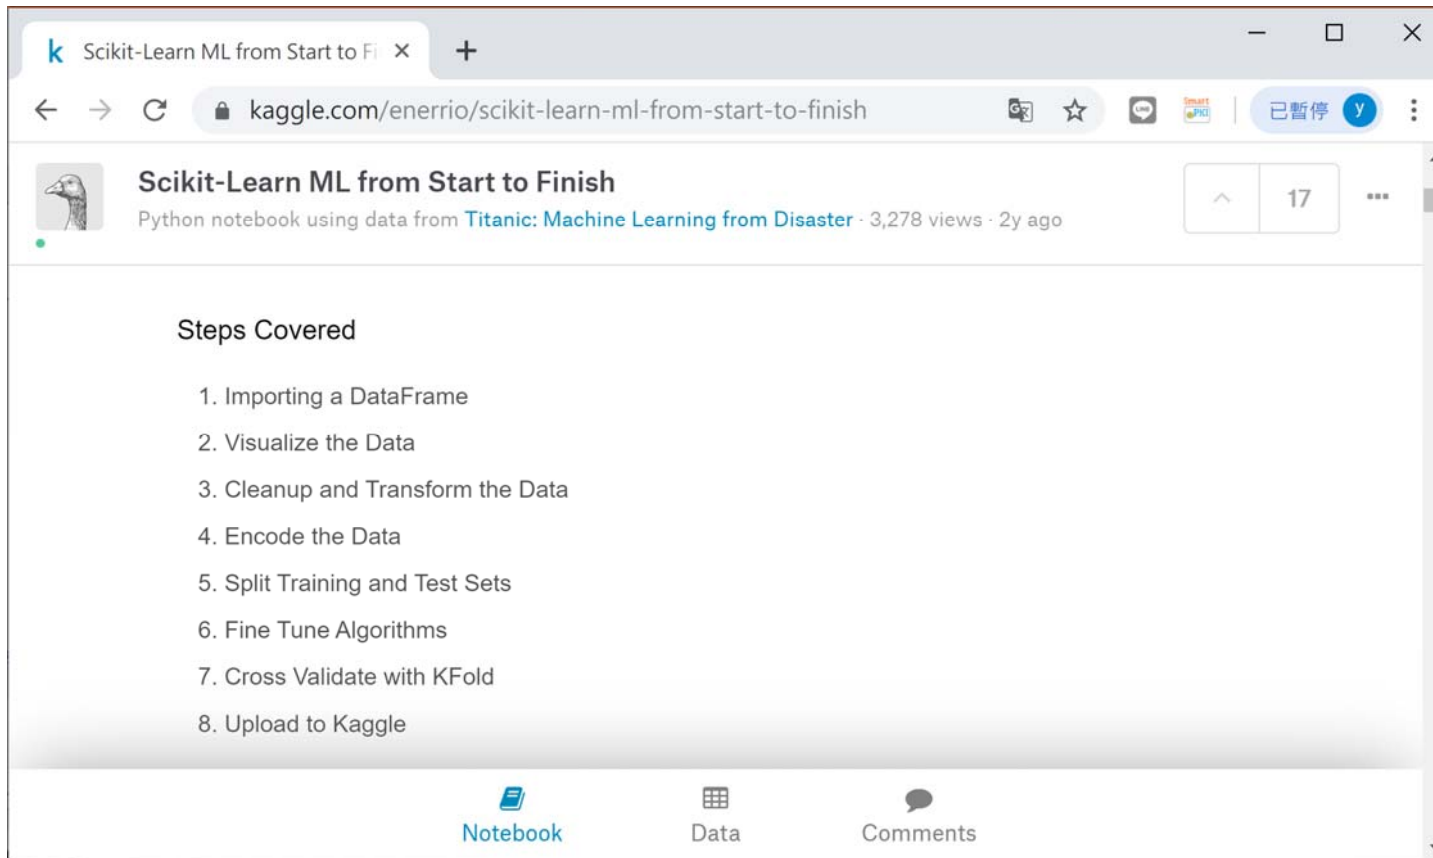
<https://www.kaggle.com/mjbahmani/the-data-scientist-s-toolbox-tutorial-1>

Data Scientist's Toolbox



<https://www.kaggle.com/mjbahmani/the-data-scientist-s-toolbox-tutorial-2>

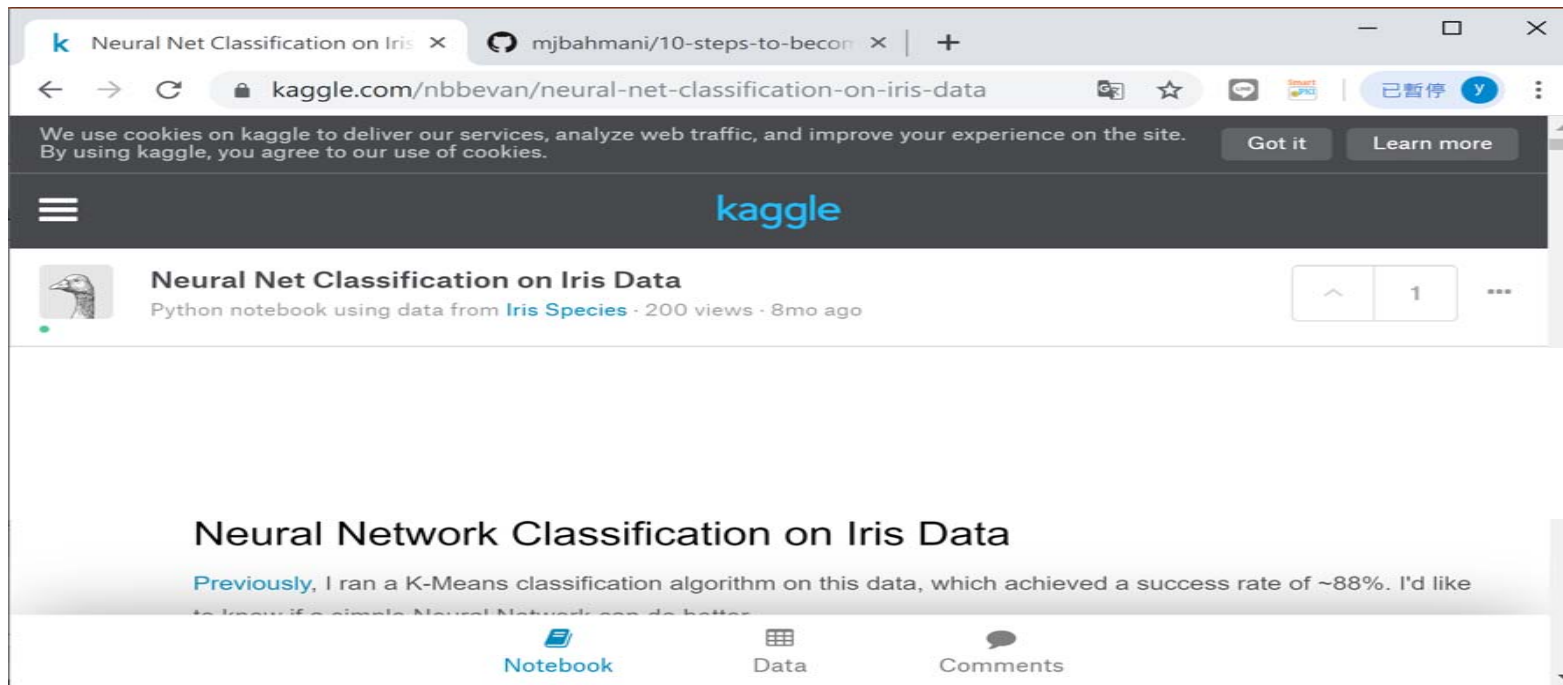
Scikit-Learn ML



<https://www.kaggle.com/enerrio/scikit-learn-ml-from-start-to-finish>

IRIS Classification

- **Neural Network**

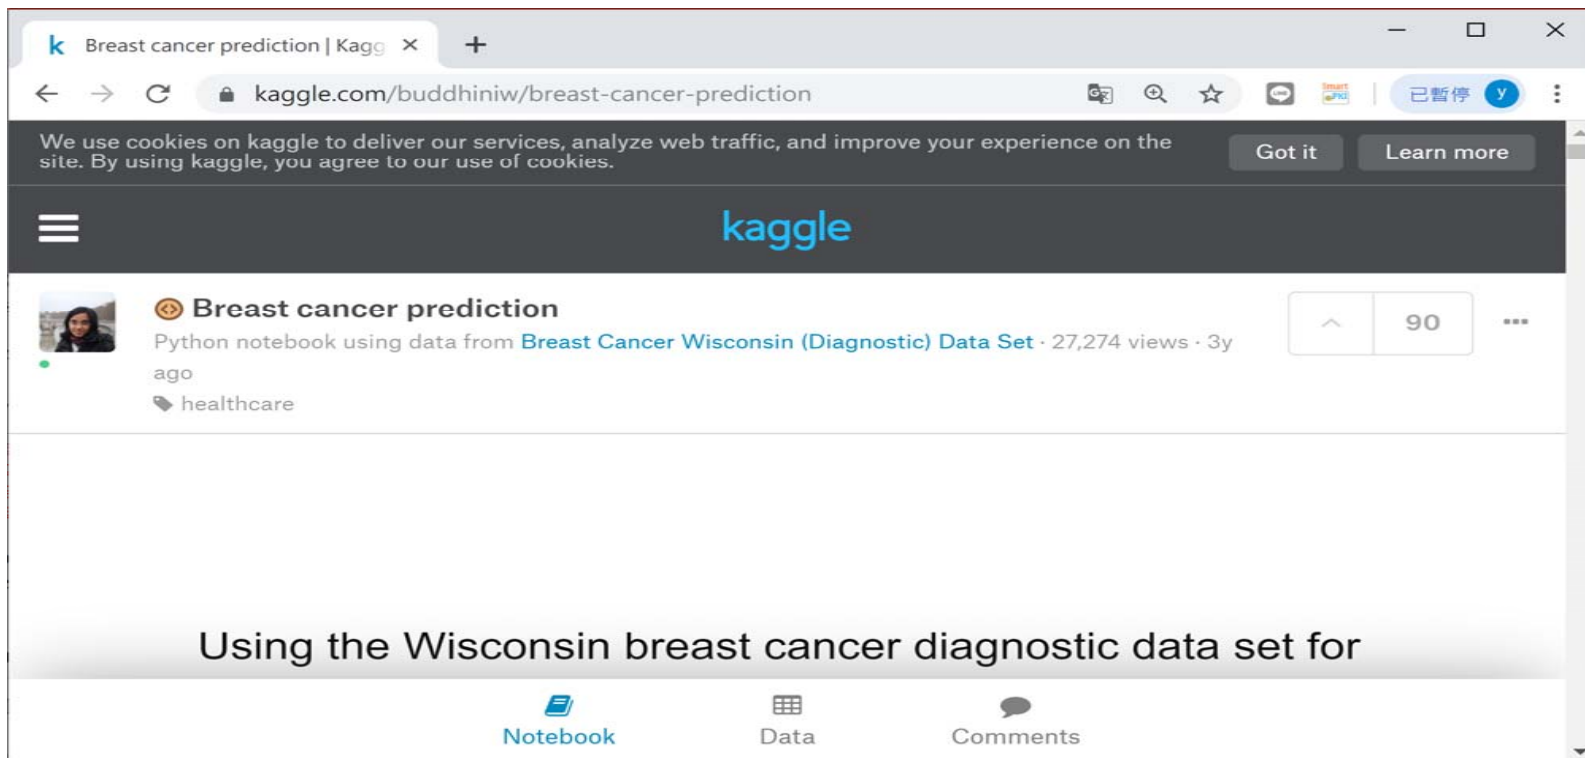


<https://www.kaggle.com/anthonyhills/classifying-species-of-iris-flowers>

<https://www.kaggle.com/mjbahmani/20-ml-algorithms-15-plot-for-beginners>


<https://www.kaggle.com/nbbevan/neural-net-classification-on-iris-data>

Breast Cancer Dataset



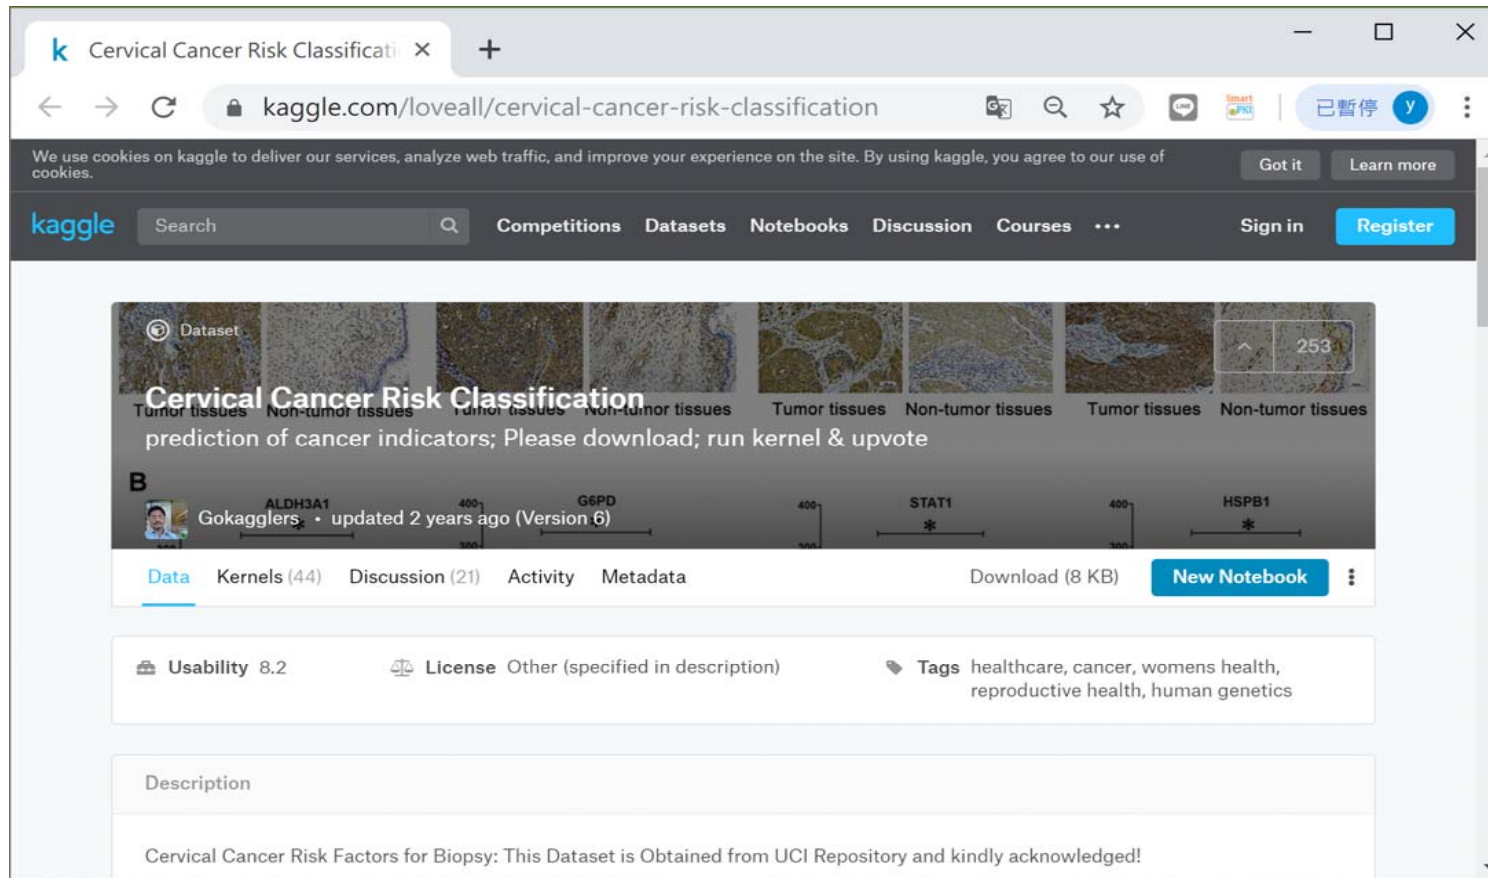
<https://www.kaggle.com/buddhiniw/breast-cancer-prediction>

Heart Disease Dataset



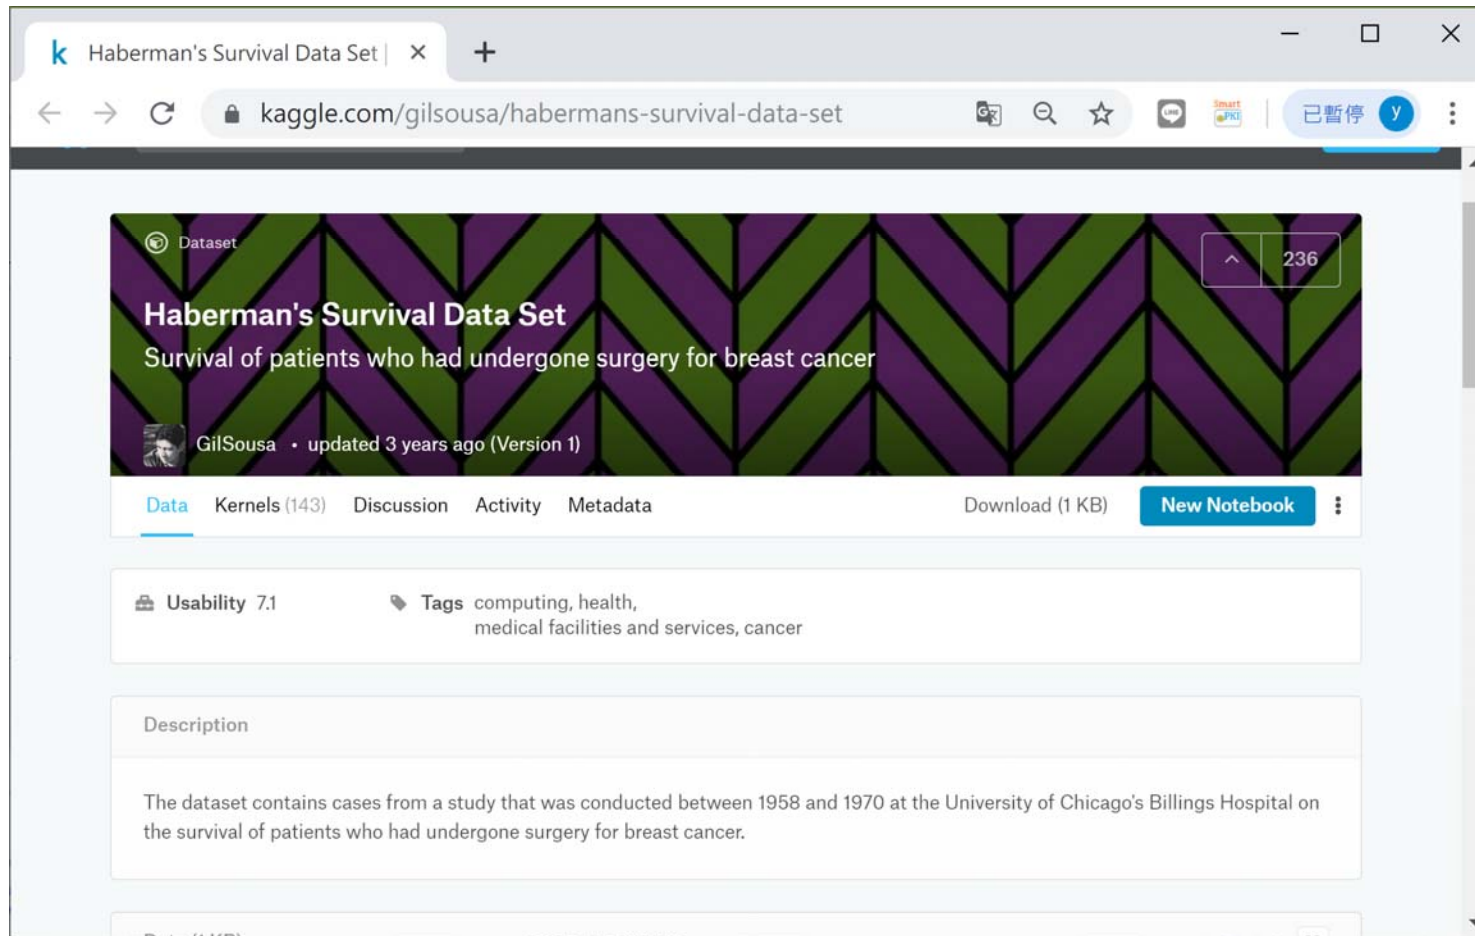
<https://www.kaggle.com/serhatmazlum/analyzing-heart-disease-s-dataset>

Cervical Cancer Risk Dataset




<https://www.kaggle.com/loveall/cervical-cancer-risk-classification>

Survival Dataset



<https://www.kaggle.com/gilsousa/habermans-survival-data-set>

Diabetes Database



<https://www.kaggle.com/buddhiniw/breast-cancer-prediction>

課後練習

- **IRIS Classification**
 - 微調模型與優化

作業二

- 使用 **Breast Cancer Dataset** 或 自行檢索的資料庫，實現**Neural Networks**。